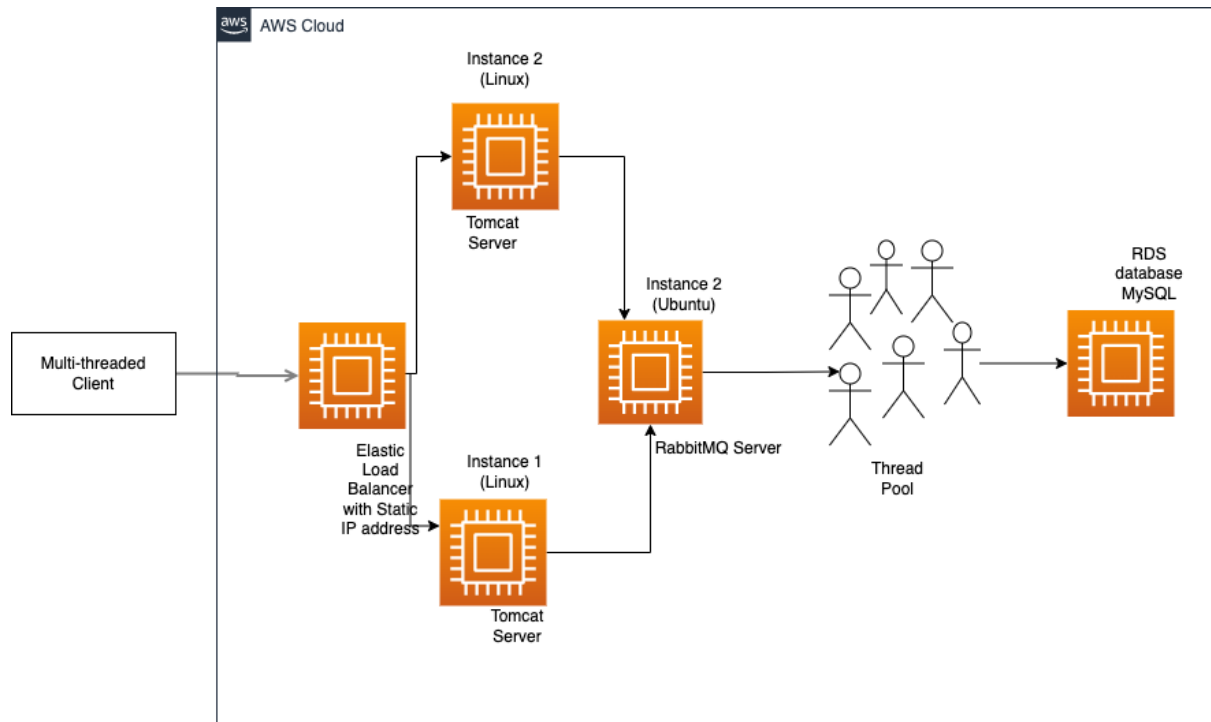


**Github URL:**

<https://github.com/NabeelHR/CS6650-Milestone-2>

**Description:**

The architecture




My architecture has the following main components:

1. A multi-threaded client (from lab 4) running locally that bombards the server with POST requests
2. A load balancer with a static IP address
3. 2 tomcat servers running on an EC2 instances that sends the appropriate messages on a rabbitMQ channel and responds to the client's queries
4. Another EC2 instance (Ubuntu version) that hosts the rabbitMQ server
5. A multi-threaded consumer also on EC2 that consumes the threads and commits the messages to a database (mySQL)

## Results and Explanations


### 64 threads

Queued messages **last minute** ?



Ready	0
Unacked	0
Total	0


Message rates **last minute** ?



Publish	0.00/s
Publisher confirm	0.00/s
Deliver (manual ack)	0.00/s


### 128 threads

Queued messages **last minute** ?



Ready	0
Unacked	0
Total	0

Message rates **last minute** ?



Publish	272/s
Publisher confirm	0.00/s
Deliver (manual ack)	0.00/s


## 256 threads

Queued messages **last minute** ?



Ready	0
Unacked	0
Total	0

Message rates **last minute** ?



Publish	278/s
Publisher confirm	0.00/s
Deliver (manual ack)	0.00/s

*All clients were run with a constant number of skiers ie 1024*

Key points:

- The queue never gets piled up
- Rate of messages incoming and messages ACKed is fairly constant across increasing client threads